

# 041-sqlite

May 18, 2022

## 4.1. Wrangling Data with SQL

```
[1]: import sqlite3

import pandas as pd
from IPython.display import VimeoVideo

[2]: VimeoVideo("665414044", h="ff34728e6a", width=600)

[2]: <IPython.lib.display.VimeoVideo at 0x7fe13eeb3490>
```

## 1 Prepare Data

### 1.1 Connect

```
[3]: VimeoVideo("665414180", h="573444d2f6", width=600)

[3]: <IPython.lib.display.VimeoVideo at 0x7fe13eeb3880>
```

**Task 4.1.1:** Run the cell below to connect to the `nepal.sqlite` database.

- What's `ipython-sql`?
- What's a Magics function?

```
[4]: %load_ext sql
      %sql sqlite:///home/jovyan/nepal.sqlite

[4]: 'Connected: @/home/jovyan/nepal.sqlite'
```

### 1.2 Explore

```
[5]: VimeoVideo("665414201", h="4f30b7a95f", width=600)

[5]: <IPython.lib.display.VimeoVideo at 0x7fe13e893d00>
```

**Task 4.1.2:** Select all rows and columns from the `sqlite_schema` table, and examine the output.

- What's a SQL database?
- What's a SQL table?

- Write a basic query in SQL.

How many tables are in the `nepal.sqlite` database? What information do they hold?

```
[6]: %%sql
SELECT *
FROM sqlite_schema

* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[6]: [('table', 'id_map', 'id_map', 2, 'CREATE TABLE "id_map" (\n"household_id"
INTEGER,\n "building_id" INTEGER,\n "vdcmun_id" INTEGER,\n "district_id"
INTEGER\n)'),
('index', 'ix_id_map_household_id', 'id_map', 3, 'CREATE INDEX
"ix_id_map_household_id"ON "id_map" ("household_id")'),
('table', 'building_structure', 'building_structure', 2032, 'CREATE TABLE
"building_structure" (\n"building_id" INTEGER,\n "count_floors_pre_eq"
INTEGER,\n "count_floors_post_eq" INTEGER,\n "age_building" IN ... (198
characters truncated) ... or_type" TEXT,\n "other_floor_type" TEXT,\n
"position" TEXT,\n "plan_configuration" TEXT,\n "condition_post_eq" TEXT,\n
"superstructure" TEXT\n)'),
('index', 'ix_building_structure_building_id', 'building_structure', 2033,
'CREATE INDEX "ix_building_structure_building_id"ON "building_structure"
("building_id")'),
('table', 'building_damage', 'building_damage', 12302, 'CREATE TABLE
"building_damage" (\n"building_id" INTEGER,\n "damage_overall_collapse" TEXT,\n
"damage_overall_leaning" TEXT,\n "damage_overall_adja ... (2923 characters
truncated) ... ndslide" INTEGER,\n "has_geotechnical_risk_rock_fall" INTEGER,\n
"has_geotechnical_risk_flood" INTEGER,\n "has_geotechnical_risk_other"
INTEGER\n)'),
('index', 'ix_building_damage_building_id', 'building_damage', 12305, 'CREATE
INDEX "ix_building_damage_building_id"ON "building_damage" ("building_id")'),
('table', 'household_demographics', 'household_demographics', 31601, 'CREATE
TABLE "household_demographics" (\n"household_id" INTEGER,\n
"gender_household_head" TEXT,\n "age_household_head" REAL,\n "caste_household"
... (8 characters truncated) ... "education_level_household_head" TEXT,\n
"income_level_household" TEXT,\n "size_household" REAL,\n
"is_bank_account_present_in_household" REAL\n)'),
('index', 'ix_household_demographics_household_id', 'household_demographics',
31602, 'CREATE INDEX "ix_household_demographics_household_id"ON
"household_demographics" ("household_id")')]
```

```
[7]: VimeoVideo("665414239", h="d7319aa0a8", width=600)
```

```
[7]: <IPython.lib.display.VimeoVideo at 0x7fe13e8c5640>
```

**Task 4.1.3:** Select the name column from the `sqlite_schema` table, showing only rows where the type is "table".

- Select a column from a table in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[8]: %%sql
SELECT name
FROM sqlite_schema
WHERE type = "table"
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[8]: [('id_map',),
      ('building_structure',),
      ('building_damage',),
      ('household_demographics',)]
```

```
[9]: VimeoVideo("665414263", h="5b7d1e875f", width=600)
```

```
[9]: <IPython.lib.display.VimeoVideo at 0x7fe13e8c5040>
```

**Task 4.1.4:** Select all columns from the `id_map` table, limiting your results to the first five rows.

- Inspect a table using a `LIMIT` clause in SQL.

How is the data organized? What type of observation does each row represent? How do you think the `household_id`, `building_id`, `vdcmun_id`, and `district_id` columns are related to each other?

```
[10]: %%sql
SELECT *
FROM id_map
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[10]: [(5601, 56, 7, 1),
      (6301, 63, 7, 1),
      (9701, 97, 7, 1),
      (9901, 99, 7, 1),
      (11501, 115, 7, 1)]
```

```
[11]: VimeoVideo("665414293", h="72fbe6b7d8", width=600)
```

```
[11]: <IPython.lib.display.VimeoVideo at 0x7fe13e8677c0>
```

**Task 4.1.5:** How many observations are in the `id_map` table? Use the `count` command to find out.

- Calculate the number of rows in a table using a `count` function in SQL.

```
[12]: %%sql
      SELECT count(*)
      FROM id_map
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[12]: [(249932,)]
```

```
[13]: VimeoVideo("665414303", h="6ba10ddf88", width=600)
```

```
[13]: <IPython.lib.display.VimeoVideo at 0x7fe13e871490>
```

**Task 4.1.6:** What districts are represented in the `id_map` table? Use the `distinct` command to determine the unique values in the `district_id` column.

- Determine the unique values in a column using a `distinct` function in SQL.

```
[14]: %%sql
      SELECT distinct(district_id)
      FROM id_map
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[14]: [(1,), (2,), (3,), (4,)]
```

```
[15]: VimeoVideo("665414313", h="adbab3e418", width=600)
```

```
[15]: <IPython.lib.display.VimeoVideo at 0x7fe13e8c5d00>
```

**Task 4.1.7:** How many buildings are there in `id_map` table? Combine the `count` and `distinct` commands to calculate the number of unique values in `building_id`.

- Calculate the number of rows in a table using a `count` function in SQL.
- Determine the unique values in a column using a `distinct` function in SQL.

```
[16]: %%sql
      SELECT count(distinct(building_id))
      FROM id_map
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[16]: [(234835,)]
```

```
[17]: VimeoVideo("665414336", h="5b595107c6", width=600)
```

```
[17]: <IPython.lib.display.VimeoVideo at 0x7fe13e871790>
```

**Task 4.1.8:** For our model, we'll focus on Gorkha (district 4). Select all columns that from `id_map`, showing only rows where the `district_id` is 4 and limiting your results to the first five rows.

- Inspect a table using a `LIMIT` clause in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[18]: %%sql
      SELECT *
      FROM id_map
      WHERE district_id = 4
      LIMIT 5

* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[18]: [(16400201, 164002, 38, 4),
      (16408101, 164081, 38, 4),
      (16408901, 164089, 38, 4),
      (16409801, 164098, 38, 4),
      (16410301, 164103, 38, 4)]
```

```
[19]: VimeoVideo("665414344", h="bdc4a50a3", width=600)
```

```
[19]: <IPython.lib.display.VimeoVideo at 0x7fe13e871b50>
```

**Task 4.1.9:** How many observations in the `id_map` table come from Gorkha? Use the `count` and `WHERE` commands together to calculate the answer.

- Calculate the number of rows in a table using a `count` function in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[20]: %%sql
      SELECT count(*)
      FROM id_map
      WHERE district_id = 4

* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[20]: [(75883,)]
```

```
[21]: VimeoVideo("665414356", h="5d2bdb3813", width=600)
```

```
[21]: <IPython.lib.display.VimeoVideo at 0x7fe13e871a60>
```

**Task 4.1.10:** How many buildings in the `id_map` table are in Gorkha? Combine the `count` and `distinct` commands to calculate the number of unique values in `building_id`, considering only rows where the `district_id` is 4.

- Calculate the number of rows in a table using a `count` function in SQL.

- Determine the unique values in a column using a `distinct` function in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[22]: %%sql
SELECT count(distinct(building_id)) AS unique_buildings_gorkha
FROM id_map
WHERE district_id = 4
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[22]: [(70836,)]
```

```
[23]: VimeoVideo("665414390", h="308ea86e4b", width=600)
```

```
[23]: <IPython.lib.display.VimeoVideo at 0x7fe13e882640>
```

**Task 4.1.11:** Select all the columns from the `building_structure` table, and limit your results to the first five rows.

- Inspect a table using a `LIMIT` clause in SQL.

What information is in this table? What does each row represent? How does it relate to the information in the `id_map` table?

```
[24]: %%sql
SELECT *
FROM building_structure
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[24]: [(1, 1, 1, 9, 288, 9, 9, 'Flat', 'Other', 'Bamboo/Timber-Light roof', 'Mud',
'Not applicable', 'Not attached', 'Rectangular', 'Damaged-Used in risk', 'Stone,
mud mortar'),
(2, 1, 1, 15, 364, 9, 9, 'Flat', 'Other', 'Bamboo/Timber-Light roof', 'Mud',
'Not applicable', 'Not attached', 'Rectangular', 'Damaged-Repaired and used',
'Stone, mud mortar'),
(3, 1, 1, 20, 384, 9, 9, 'Flat', 'Other', 'Bamboo/Timber-Light roof', 'Mud',
'Not applicable', 'Not attached', 'Rectangular', 'Damaged-Repaired and used',
'Stone, mud mortar'),
(4, 1, 1, 20, 312, 9, 9, 'Flat', 'Other', 'Bamboo/Timber-Light roof', 'Mud',
'Not applicable', 'Not attached', 'Rectangular', 'Damaged-Repaired and used',
'Stone, mud mortar'),
(5, 1, 1, 30, 308, 9, 9, 'Flat', 'Other', 'Bamboo/Timber-Light roof', 'Mud',
'Not applicable', 'Not attached', 'Rectangular', 'Damaged-Repaired and used',
'Stone, mud mortar')]
```

```
[25]: VimeoVideo("665414402", h="64875c7779", width=600)
```

```
[25]: <IPython.lib.display.VimeoVideo at 0x7fe13e882a60>
```

**Task 4.1.12:** How many building are there in the `building_structure` table? Use the `count` command to find out.

- Calculate the number of rows in a table using a `count` function in SQL.

```
[26]: %%sql
      SELECT count(*)
      FROM building_structure
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[26]: [(234835,)]
```

```
[27]: VimeoVideo("665414414", h="202f83f3cb", width=600)
```

```
[27]: <IPython.lib.display.VimeoVideo at 0x7fe13e80c2b0>
```

**Task 4.1.13:** There are over 200,000 buildings in the `building_structure` table, but how can we retrieve only buildings that are in Gorkha? Use the `JOIN` command to join the `id_map` and `building_structure` tables, showing only buildings where `district_id` is 4 and limiting your results to the first five rows of the new table.

- Create an alias for a column or table using the `AS` command in SQL.
- Merge two tables using a `JOIN` clause in SQL.
- Inspect a table using a `LIMIT` clause in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[33]: %%sql
      SELECT *
      FROM id_map AS i
      JOIN building_structure AS s ON i.building_id = s.building_id
      WHERE district_id = 4
      LIMIT 5
```

```
# Above code is a example of Left Join
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[33]: [(16400201, 164002, 38, 4, 164002, 3, 3, 20, 560, 18, 18, 'Flat', 'Mud mortar-
      Stone/Brick', 'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not
      attached', 'Rectangular', 'Damaged-Repaired and used', 'Stone, mud mortar'),
      (16408101, 164081, 38, 4, 164081, 2, 2, 21, 200, 12, 12, 'Flat', 'Mud mortar-
      Stone/Brick', 'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not
```

```
attached', 'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'),
(16408901, 164089, 38, 4, 164089, 3, 3, 18, 315, 20, 20, 'Flat', 'Mud mortar-
Stone/Brick', 'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not
attached', 'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'),
(16409801, 164098, 38, 4, 164098, 2, 2, 45, 290, 13, 13, 'Flat', 'Mud mortar-
Stone/Brick', 'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not
attached', 'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'),
(16410301, 164103, 38, 4, 164103, 2, 2, 21, 230, 13, 13, 'Flat', 'Mud mortar-
Stone/Brick', 'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not
attached', 'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar')]
```

In the table we just made, each row represents a unique household in Gorkha. How can we create a table where each row represents a unique building?

```
[29]: VimeoVideo("665414450", h="0fcb4dc3fa", width=600)
```

```
[29]: <IPython.lib.display.VimeoVideo at 0x7fe13e8c5f10>
```

**Task 4.1.14:** Use the `distinct` command to create a column with all unique building IDs in the `id_map` table. JOIN this column with all the columns from the `building_structure` table, showing only buildings where `district_id` is 4 and limiting your results to the first five rows of the new table.

- Create an alias for a column or table using the `AS` command in SQL.
- Determine the unique values in a column using a `distinct` function in SQL.
- Merge two tables using a `JOIN` clause in SQL.
- Inspect a table using a `LIMIT` clause in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[35]: %%sql
SELECT distinct(i.building_id),
        s.*
FROM id_map AS i
JOIN building_structure AS s ON i.building_id = s.building_id
WHERE district_id = 4
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[35]: [(164002, 164002, 3, 3, 20, 560, 18, 18, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Repaired and used', 'Stone, mud mortar'),
(164081, 164081, 2, 2, 21, 200, 12, 12, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'),
(164089, 164089, 3, 3, 18, 315, 20, 20, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
```



```
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'),
(164098, 164098, 2, 2, 45, 290, 13, 13, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'),
(164103, 164103, 2, 2, 21, 230, 13, 13, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar'))]
```

We've combined the `id_map` and `building_structure` tables to create a table with all the buildings in Gorkha, but the final piece of data needed for our model, the damage that each building sustained in the earthquake, is in the `building_damage` table.

```
[36]: VimeoVideo("665414466", h="37dde03c93", width=600)
```

```
[36]: <IPython.lib.display.VimeoVideo at 0x7fe13e871670>
```

**Task 4.1.15:** How can combine all three tables? Using the query you created in the last task as a foundation, include the `damage_grade` column to your table by adding a second JOIN for the `building_damage` table. Be sure to limit your results to the first five rows of the new table.

- Create an alias for a column or table using the `AS` command in SQL.
- Determine the unique values in a column using a `distinct` function in SQL.
- Merge two tables using a `JOIN` clause in SQL.
- Inspect a table using a `LIMIT` clause in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[40]: %%sql
SELECT distinct(i.building_id) AS b_id,
        s.*,
        d.damage_grade
FROM id_map AS i
JOIN building_structure AS s ON i.building_id = s.building_id
JOIN building_damage AS d ON i.building_id = d.building_id
WHERE district_id = 4
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[40]: [(164002, 164002, 3, 3, 20, 560, 18, 18, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Repaired and used', 'Stone, mud mortar', 'Grade 2'),
(164081, 164081, 2, 2, 21, 200, 12, 12, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar', 'Grade 2'),
(164089, 164089, 3, 3, 18, 315, 20, 20, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar', 'Grade 2'),
```

```
(164098, 164098, 2, 2, 45, 290, 13, 13, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar', 'Grade 3'),
(164103, 164103, 2, 2, 21, 230, 13, 13, 'Flat', 'Mud mortar-Stone/Brick',
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',
'Rectangular', 'Damaged-Used in risk', 'Stone, mud mortar', 'Grade 3')]
```

### 1.3 Import

```
[41]: VimeoVideo("665414492", h="9392e1a66e", width=600)
```

```
[41]: <IPython.lib.display.VimeoVideo at 0x7fe13e8671c0>
```

**Task 4.1.16:** Use the `connect` method from the `sqlite3` library to connect to the database. Remember that the database is located at `"/home/jovyan/nepal.sqlite"`.

- Open a connection to a SQL database using `sqlite3`.

```
[42]: conn = sqlite3.connect("/home/jovyan/nepal.sqlite")
```

```
[43]: VimeoVideo("665414501", h="812d482c73", width=600)
```

```
[43]: <IPython.lib.display.VimeoVideo at 0x7fe13e867e80>
```

**Task 4.1.17:** Put your last SQL query into a string and assign it to the variable `query`.

```
[44]: query = """
SELECT distinct(i.building_id) AS b_id,
        s.*,
        d.damage_grade
FROM id_map AS i
JOIN building_structure AS s ON i.building_id = s.building_id
JOIN building_damage AS d ON i.building_id = d.building_id
WHERE district_id = 4
LIMIT 5

"""
print(query)
```

```
SELECT distinct(i.building_id) AS b_id,
        s.*,
        d.damage_grade
FROM id_map AS i
JOIN building_structure AS s ON i.building_id = s.building_id
JOIN building_damage AS d ON i.building_id = d.building_id
WHERE district_id = 4
LIMIT 5
```

```
[45]: VimeoVideo("665414513", h="c6a81b49ad", width=600)
```

```
[45]: <IPython.lib.display.VimeoVideo at 0x7fe13e867fa0>
```

**Task 4.1.18:** Use the `read_sql` from the pandas library to create a DataFrame from your query. Be sure that the `building_id` is set as your index column.

- Read SQL query into a DataFrame using pandas.

Tip: Your table might have two `building_id` columns, and that will make it hard to set it as the index column for your DataFrame. If you face this problem, add an alias for one of the `building_id` columns in your query using AS.

```
[51]: df = pd.read_sql(query, conn, index_col="b_id")

df.head()
```

```
[51]:
```

	building_id	count_floors_pre_eq	count_floors_post_eq	age_building	\
b_id					
164002	164002	3	3	20	
164081	164081	2	2	21	
164089	164089	3	3	18	
164098	164098	2	2	45	
164103	164103	2	2	21	

	plinth_area_sq_ft	height_ft_pre_eq	height_ft_post_eq	\
b_id				
164002	560	18	18	
164081	200	12	12	
164089	315	20	20	
164098	290	13	13	
164103	230	13	13	

	land_surface_condition	foundation_type	\
b_id			
164002	Flat	Mud mortar-Stone/Brick	
164081	Flat	Mud mortar-Stone/Brick	
164089	Flat	Mud mortar-Stone/Brick	
164098	Flat	Mud mortar-Stone/Brick	
164103	Flat	Mud mortar-Stone/Brick	

	roof_type	ground_floor_type	other_floor_type	\
b_id				
164002	Bamboo/Timber-Light roof	Mud	Timber/Bamboo-Mud	
164081	Bamboo/Timber-Light roof	Mud	Timber/Bamboo-Mud	
164089	Bamboo/Timber-Light roof	Mud	Timber/Bamboo-Mud	

164098	Bamboo/Timber-Light roof	Mud	Timber/Bamboo-Mud
164103	Bamboo/Timber-Light roof	Mud	Timber/Bamboo-Mud

	position	plan_configuration	condition_post_eq \
b_id			
164002	Not attached	Rectangular	Damaged-Repaired and used
164081	Not attached	Rectangular	Damaged-Used in risk
164089	Not attached	Rectangular	Damaged-Used in risk
164098	Not attached	Rectangular	Damaged-Used in risk
164103	Not attached	Rectangular	Damaged-Used in risk

	superstructure	damage_grade
b_id		
164002	Stone, mud mortar	Grade 2
164081	Stone, mud mortar	Grade 2
164089	Stone, mud mortar	Grade 2
164098	Stone, mud mortar	Grade 3
164103	Stone, mud mortar	Grade 3

```
[50]: # Check your work
assert df.shape[0] == 70836, f"`df` should have 70,836 rows, not {df.shape[0]}."
assert (
    df.shape[1] > 14
), "`df` seems to be missing columns. Does your query combine the `id_map`,
↳ `building_structure`, and `building_damage` tables?"
assert (
    "damage_grade" in df.columns
), "`df` is missing the target column, `damage_grade`."
```

```
-----
AssertionError                                Traceback (most recent call last)
Input In [50], in <cell line: 2>()
      1 # Check your work
----> 2 assert df.shape[0] == 70836, f"`df` should have 70,836 rows, not {df.
↳ shape[0]}."
      3 assert (
      4     df.shape[1] > 14
      5 ), "`df` seems to be missing columns. Does your query combine the
↳ `id_map`, `building_structure`, and `building_damage` tables?"
      6 assert (
      7     "damage_grade" in df.columns
      8 ), "`df` is missing the target column, `damage_grade`."

AssertionError: `df` should have 70,836 rows, not 5.
```

Copyright © 2022 WorldQuant University. This content is licensed solely for personal use. Redistribution or publication of this material is strictly prohibited.