

June 7, 2022

4.4. Beyond the Model: Data Ethics

```
[2]: import sqlite3
import warnings

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from category_encoders import OneHotEncoder
from IPython.display import VimeoVideo
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.utils.validation import check_is_fitted

warnings.simplefilter(action="ignore", category=FutureWarning)
```

```
[3]: VimeoVideo("665414155", h="c8a3e81a05", width=600)
```

```
[3]: <IPython.lib.display.VimeoVideo at 0x7f052c45eee0>
```

1 Prepare Data

Task 4.4.1: Run the cell below to connect to the `nepal.sqlite` database.

- What's `ipython-sql`?
- What's a `Magics` function?

```
[4]: %load_ext sql
%sql sqlite:///home/jovyan/nepal.sqlite
```

```
[4]: 'Connected: @/home/jovyan/nepal.sqlite'
```

```
[5]: VimeoVideo("665415362", h="f677c48c46", width=600)
```

```
[5]: <IPython.lib.display.VimeoVideo at 0x7f04794adeb0>
```

Task 4.4.2: Select all columns from the `household_demographics` table, limiting your results to the first five rows.

- Write a basic query in SQL.
- Inspect a table using a `LIMIT` clause in SQL.

```
[6]: %%sql
SELECT *
FROM household_demographics
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[6]: [(101, 'Male', 31.0, 'Rai', 'Illiterate', 'Rs. 10 thousand', 3.0, 0.0),
      (201, 'Female', 62.0, 'Rai', 'Illiterate', 'Rs. 10 thousand', 6.0, 0.0),
      (301, 'Male', 51.0, 'Gharti/Bhujel', 'Illiterate', 'Rs. 10 thousand', 13.0,
      0.0),
      (401, 'Male', 48.0, 'Gharti/Bhujel', 'Illiterate', 'Rs. 10 thousand', 5.0,
      0.0),
      (501, 'Male', 70.0, 'Gharti/Bhujel', 'Illiterate', 'Rs. 10 thousand', 8.0,
      0.0)]
```

Task 4.4.3: How many observations are in the `household_demographics` table? Use the `count` command to find out.

- Calculate the number of rows in a table using a `count` function in SQL.

```
[7]: %%sql
SELECT count(*)
FROM household_demographics
```

```
* sqlite:///home/jovyan/nepal.sqlite
Done.
```

```
[7]: [(249932,)]
```

```
[8]: VimeoVideo("665415378", h="aa2b99493e", width=600)
```

```
[8]: <IPython.lib.display.VimeoVideo at 0x7f04794d5d00>
```

Task 4.4.4: Select all columns from the `id_map` table, limiting your results to the first five rows.

- Inspect a table using a `LIMIT` clause in SQL.

What columns does it have in common with `household_demographics` that we can use to join them?

```
[9]: %%sql
SELECT *
FROM id_map
```

```
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite  
Done.
```

```
[9]: [(5601, 56, 7, 1),  
      (6301, 63, 7, 1),  
      (9701, 97, 7, 1),  
      (9901, 99, 7, 1),  
      (11501, 115, 7, 1)]
```

```
[10]: VimeoVideo("665415406", h="46a990c8f7", width=600)
```

```
[10]: <IPython.lib.display.VimeoVideo at 0x7f04794adb80>
```

Task 4.4.5: Create a table with all the columns from `household_demographics`, all the columns from `building_structure`, the `vdcmun_id` column from `id_map`, and the `damage_grade` column from `building_damage`. Your results should show only rows where the `district_id` is 5 and limit your results to the first five rows.

- Create an alias for a column or table using the `AS` command in SQL.
- Determine the unique values in a column using a `DISTINCT` function in SQL.
- Merge two tables using a `JOIN` clause in SQL.
- Inspect a table using a `LIMIT` clause in SQL.
- Subset a table using a `WHERE` clause in SQL.

```
[11]: %%sql  
SELECT h.*,  
       s.*,  
       i.vdcmun_id,  
       d.damage_grade  
FROM household_demographics AS h  
JOIN id_map AS i ON i.household_id = h.household_id  
JOIN building_structure AS s ON i.building_id = s.building_id  
JOIN building_damage AS d ON i.building_id = d.building_id  
WHERE district_id = 4  
LIMIT 5
```

```
* sqlite:///home/jovyan/nepal.sqlite  
Done.
```

```
[11]: [(16400201, 'Female', 46.0, 'Chhetree', 'Class 5', 'Rs. 10-20 thousand', 4.0,  
1.0, 164002, 3, 3, 20, 560, 18, 18, 'Flat', 'Mud mortar-Stone/Brick',  
'Bamboo/Timber-Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached',  
'Rectangular', 'Damaged-Repaired and used', 'Stone, mud mortar', 38, 'Grade 2'),  
(16408101, 'Male', 66.0, 'Chhetree', 'Illiterate', 'Rs. 10 thousand', 5.0, 0.0,  
164081, 2, 2, 21, 200, 12, 12, 'Flat', 'Mud mortar-Stone/Brick', 'Bamboo/Timber-  
Light roof', 'Mud', 'Timber/Bamboo-Mud', 'Not attached', 'Rectangular',
```

```
'Damaged-Used in risk', 'Stone, mud mortar', 38, 'Grade 2'),
(16408901, 'Male', 54.0, 'Magar', 'Class 4', 'Rs. 10 thousand', 5.0, 1.0,
164089, 3, 3, 18, 315, 20, 20, 'Flat', 'Mud mortar-Stone/Brick', 'Bamboo/Timber-
Light roof', 'Mud', 'TImber/Bamboo-Mud', 'Not attached', 'Rectangular',
'Damaged-Used in risk', 'Stone, mud mortar', 38, 'Grade 2'),
(16409801, 'Male', 36.0, 'Chhetree', 'Class 5', 'Rs. 10 thousand', 6.0, 1.0,
164098, 2, 2, 45, 290, 13, 13, 'Flat', 'Mud mortar-Stone/Brick', 'Bamboo/Timber-
Light roof', 'Mud', 'TImber/Bamboo-Mud', 'Not attached', 'Rectangular',
'Damaged-Used in risk', 'Stone, mud mortar', 38, 'Grade 3'),
(16410301, 'Female', 39.0, 'Chhetree', 'Class 4', 'Rs. 10 thousand', 3.0, 0.0,
164103, 2, 2, 21, 230, 13, 13, 'Flat', 'Mud mortar-Stone/Brick', 'Bamboo/Timber-
Light roof', 'Mud', 'TImber/Bamboo-Mud', 'Not attached', 'Rectangular',
'Damaged-Used in risk', 'Stone, mud mortar', 38, 'Grade 3')]
```

1.1 Import

```
[12]: def wrangle(db_path):
    # Connect to database
    conn = sqlite3.connect(db_path)

    # Construct query
    query = """
    SELECT h.*,
           s.*,
           i.vdcmun_id,
           d.damage_grade
    FROM household_demographics AS h
    JOIN id_map AS i ON i.household_id = h.household_id
    JOIN building_structure AS s ON i.building_id = s.building_id
    JOIN building_damage AS d ON i.building_id = d.building_id
    WHERE district_id = 4
    """

    # Read query results into DataFrame
    df = pd.read_sql(query, conn, index_col="household_id")

    # Identify leaky columns
    drop_cols = [col for col in df.columns if "post_eq" in col]

    # Add high-cardinality / redundant column
    drop_cols.append("building_id")

    # Create binary target column
    df["damage_grade"] = df["damage_grade"].str[-1].astype(int)
    df["severe_damage"] = (df["damage_grade"] > 3).astype(int)

    # Drop old target
```

```

drop_cols.append("damage_grade")

# Drop multicollinearity column
drop_cols.append("count_floors_pre_eq")

#Group caste columns
top_10 = df["caste_household"].value_counts().head(10).index
df["caste_household"] = df["caste_household"].apply(
    lambda c: c if c in top_10 else "Other"
)

# Drop columns
df.drop(columns=drop_cols, inplace=True)

return df

```

```
[13]: VimeoVideo("665415443", h="ca27a7ebfc", width=600)
```

```
[13]: <IPython.lib.display.VimeoVideo at 0x7f04794d53a0>
```

Task 4.4.6: Add the query you created in the previous task to the `wrangle` function above. Then import your data by running the cell below. The path to the database is `"/home/jovyan/nepal.sqlite"`.

- Read SQL query into a DataFrame using pandas.
- Write a function in Python.

```
[14]: df = wrangle("/home/jovyan/nepal.sqlite")
df.head(10)
```

```
[14]:
```

household_id	gender_household_head	age_household_head	caste_household
16400201	Female	46.0	Chhetree
16408101	Male	66.0	Chhetree
16408901	Male	54.0	Magar
16409801	Male	36.0	Chhetree
16410301	Female	39.0	Chhetree
16418601	Female	50.0	Sarki
16420401	Female	48.0	Magar
16420501	Female	55.0	Magar
16421101	Male	44.0	Magar
16422001	Male	46.0	Magar

household_id	education_level_household_head	income_level_household
16400201	Class 5	Rs. 10-20 thousand
16408101	Illiterate	Rs. 10 thousand
16408901	Class 4	Rs. 10 thousand

16409801	Class 5	Rs. 10 thousand
16410301	Class 4	Rs. 10 thousand
16418601	Illiterate	Rs. 10 thousand
16420401	Intermediate or equivalent	Rs. 10-20 thousand
16420501	Intermediate or equivalent	Rs. 10-20 thousand
16421101	Class 10	Rs. 10 thousand
16422001	Class 6	Rs. 10 thousand

household_id	size_household	is_bank_account_present_in_household	\
16400201	4.0	1.0	
16408101	5.0	0.0	
16408901	5.0	1.0	
16409801	6.0	1.0	
16410301	3.0	0.0	
16418601	5.0	0.0	
16420401	4.0	1.0	
16420501	4.0	1.0	
16421101	5.0	0.0	
16422001	6.0	0.0	

household_id	age_building	plinth_area_sq_ft	height_ft_pre_eq	\
16400201	20	560	18	
16408101	21	200	12	
16408901	18	315	20	
16409801	45	290	13	
16410301	21	230	13	
16418601	40	250	13	
16420401	20	350	13	
16420501	45	400	13	
16421101	40	250	21	
16422001	4	300	12	

household_id	land_surface_condition	foundation_type	\
16400201	Flat	Mud mortar-Stone/Brick	
16408101	Flat	Mud mortar-Stone/Brick	
16408901	Flat	Mud mortar-Stone/Brick	
16409801	Flat	Mud mortar-Stone/Brick	
16410301	Flat	Mud mortar-Stone/Brick	
16418601	Flat	Mud mortar-Stone/Brick	
16420401	Flat	Mud mortar-Stone/Brick	
16420501	Flat	Mud mortar-Stone/Brick	
16421101	Moderate slope	Mud mortar-Stone/Brick	
16422001	Flat	Mud mortar-Stone/Brick	

household_id	roof_type	ground_floor_type	other_floor_type	\
16400201	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16408101	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16408901	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16409801	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16410301	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16418601	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16420401	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16420501	Bamboo/Timber-Heavy roof		Mud	Timber/Bamboo-Mud
16421101	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud
16422001	Bamboo/Timber-Light roof		Mud	Timber/Bamboo-Mud

household_id	position	plan_configuration	superstructure	\
16400201	Not attached	Rectangular	Stone, mud mortar	
16408101	Not attached	Rectangular	Stone, mud mortar	
16408901	Not attached	Rectangular	Stone, mud mortar	
16409801	Not attached	Rectangular	Stone, mud mortar	
16410301	Not attached	Rectangular	Stone, mud mortar	
16418601	Not attached	Rectangular	Stone, mud mortar	
16420401	Not attached	Rectangular	Stone, mud mortar	
16420501	Not attached	Rectangular	Stone, mud mortar	
16421101	Attached-2 side	Rectangular	Stone, mud mortar	
16422001	Not attached	Rectangular	Stone, mud mortar	

household_id	vdcmun_id	severe_damage
16400201	38	0
16408101	38	0
16408901	38	0
16409801	38	0
16410301	38	0
16418601	38	1
16420401	38	1
16420501	38	1
16421101	38	1
16422001	38	1

```
[15]: # Check your work
assert df.shape == (75883, 20), f"`df` should have shape (75883, 20), not {df.
↳shape}"
```

1.2 Explore

```
[16]: VimeoVideo("665415463", h="86c306199f", width=600)
```

```
[16]: <IPython.lib.display.VimeoVideo at 0x7f04794d50d0>
```

Task 4.4.7: Combine the `select_dtypes` and `nunique` methods to see if there are any high- or low-cardinality categorical features in the dataset.

- What are high- and low-cardinality features?
- Determine the unique values in a column using pandas.
- Subset a DataFrame's columns based on the column data types in pandas.

```
[17]: # Check for high- and low-cardinality categorical features
df.select_dtypes("object").nunique()
```

```
[17]: gender_household_head      2
caste_household              11
education_level_household_head 19
income_level_household        5
land_surface_condition        3
foundation_type               5
roof_type                    3
ground_floor_type             5
other_floor_type              4
position                     4
plan_configuration            10
superstructure                11
dtype: int64
```

```
[18]: VimeoVideo("665415472", h="1142d69e4a", width=600)
```

```
[18]: <IPython.lib.display.VimeoVideo at 0x7f04794d5970>
```

```
[19]: # top_10 = df["caste_household"].value_counts().head(10).index
# df["caste_household"].apply(lambda c: c if c in top_10 else "Other").
↪value_counts()
```

Task 4.4.8: Add to your `wrangle` function so that the "caste_household" contains only the 10 largest caste groups. For the rows that are not in those groups, "caste_household" should be changed to "Other".

- Determine the unique values in a column using pandas.
- Combine multiple categories in a Series using pandas.

```
[20]: df["caste_household"].nunique()
```

```
[20]: 11
```



```
[21]: # Check your work
assert (
    df["caste_household"].nunique() == 11
), f"The `'caste_household`' column should only have 11 unique values, not {df['caste_household'].nunique()}."
```

1.3 Split

```
[22]: VimeoVideo("665415515", h="defc252edd", width=600)
```

```
[22]: <IPython.lib.display.VimeoVideo at 0x7f04794d5af0>
```

Task 4.4.9: Create your feature matrix X and target vector y . Since our model will only consider building and household data, X should not include the municipality column "vdcmun_id". Your target is "severe_damage".

```
[23]: target = "severe_damage"
X = df.drop(columns=[target, "vdcmun_id"])
y = df[target]
```

```
[24]: # Check your work
assert X.shape == (75883, 18), f"The shape of `X` should be (75883, 18), not {X.shape}."
assert "vdcmun_id" not in X.columns, "There should be no `'vdcmun_id`' column in `X`."
assert y.shape == (75883,), f"The shape of `y` should be (75883,), not {y.shape}."
```

Task 4.4.10: Divide your data (X and y) into training and test sets using a randomized train-test split. Your test set should be 20% of your total data. Be sure to set a `random_state` for reproducibility.

```
[25]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
[26]: # Check your work
assert X_train.shape == (
    60706,
    18,
), f"The shape of `X_train` should be (60706, 18), not {X_train.shape}."
assert y_train.shape == (
    60706,
), f"The shape of `y_train` should be (60706,), not {y_train.shape}."
assert X_test.shape == (
    15177,
    18,
```

```
), f"The shape of `X_test` should be (15177, 18), not {X_test.shape}."
assert y_test.shape == (
    15177,
), f"The shape of `y_test` should be (15177,), not {y_test.shape}."
```

2 Build Model

2.1 Baseline

Task 4.4.11: Calculate the baseline accuracy score for your model.

- What's accuracy score?
- Aggregate data in a Series using `value_counts` in pandas.

```
[27]: acc_baseline = y_train.value_counts(normalize=True).max()
print("Baseline Accuracy:", round(acc_baseline, 2))
```

Baseline Accuracy: 0.63

2.2 Iterate

Task 4.4.12: Create a Pipeline called `model_lr`. It should have an `OneHotEncoder` transformer and a `LogisticRegression` predictor. Be sure you set the `use_cat_names` argument for your transformer to `True`.

- What's logistic regression?
- What's one-hot encoding?
- Create a pipeline in scikit-learn.
- Fit a model to training data in scikit-learn.

```
[28]: model_lr = make_pipeline(
    OneHotEncoder(use_cat_names=True),
    LogisticRegression(max_iter=3000)
)
model_lr.fit(X_train, y_train)
```

```
[28]: Pipeline(steps=[('onehotencoder',
    OneHotEncoder(cols=['gender_household_head', 'caste_household',
    'education_level_household_head',
    'income_level_household',
    'land_surface_condition',
    'foundation_type', 'roof_type',
    'ground_floor_type', 'other_floor_type',
    'position', 'plan_configuration',
    'superstructure'],
    use_cat_names=True)),
    ('logisticregression', LogisticRegression(max_iter=3000))])
```

```
[29]: # Check your work
assert isinstance(
    model_lr, Pipeline
), f"`model_lr` should be a Pipeline, not type {type(model_lr)}."
assert isinstance(
    model_lr[0], OneHotEncoder
), f"The first step in your Pipeline should be a OneHotEncoder, not type_
↳{type(model_lr[0])}."
assert isinstance(
    model_lr[-1], LogisticRegression
), f"The last step in your Pipeline should be LogisticRegression, not type_
↳{type(model_lr[-1])}."
check_is_fitted(model_lr)
```

2.3 Evaluate

Task 4.4.13: Calculate the training and test accuracy scores for `model_lr`.

- Calculate the accuracy score for a model in scikit-learn.
- Generate predictions using a trained model in scikit-learn.

```
[30]: acc_train = accuracy_score(y_train, model_lr.predict(X_train))
acc_test = model_lr.score(X_test, y_test)

print("LR Training Accuracy:", acc_train)
print("LR Validation Accuracy:", acc_test)
```

LR Training Accuracy: 0.7181497710275755

LR Validation Accuracy: 0.7220135731699282

3 Communicate

```
[31]: VimeoVideo("665415532", h="00440f76a9", width=600)
```

```
[31]: <IPython.lib.display.VimeoVideo at 0x7f0471ba4f70>
```

Task 4.4.14: First, extract the feature names and importances from your model. Then create a pandas Series named `feat_imp`, where the index is `features` and the values are your the exponential of the importances.

- What's a bar chart?
- Access an object in a pipeline in scikit-learn.
- Create a Series in pandas.

```
[32]: features = model_lr.named_steps["onehotencoder"].get_feature_names()
importances = model_lr.named_steps["logisticregression"].coef_[0]
feat_imp = pd.Series(np.exp(importances), index=features).sort_values()
feat_imp.head()
```

```
[32]: superstructure_Brick, cement mortar    0.320384
      foundation_type_RC                    0.352191
      roof_type_RCC/RB/RBC                  0.413963
      ground_floor_type_RC                  0.535611
      caste_household_Kumal                 0.540619
      dtype: float64
```

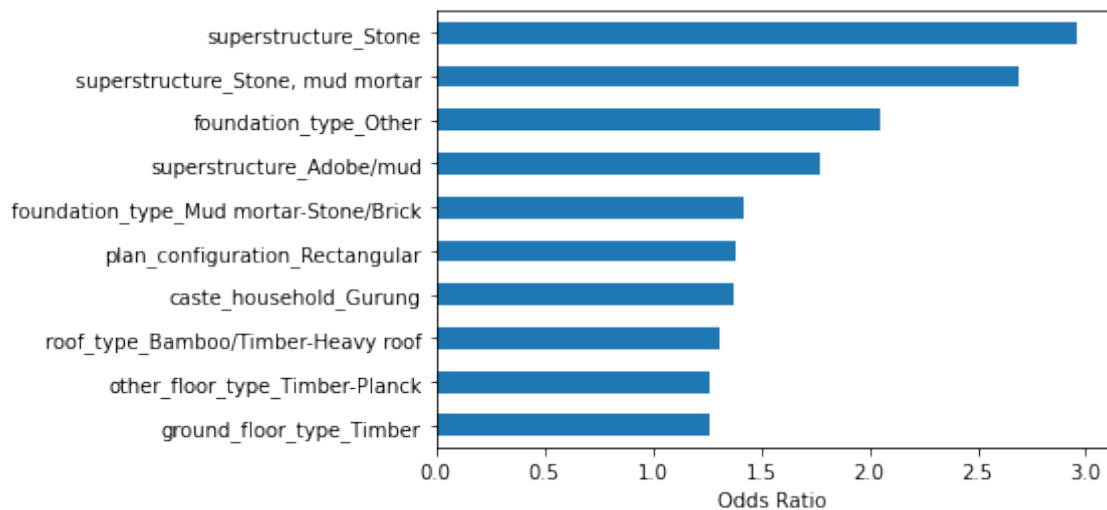
```
[33]: VimeoVideo("665415552", h="5b2383ccf8", width=600)
```

```
[33]: <IPython.lib.display.VimeoVideo at 0x7f0471ba4d00>
```

Task 4.4.15: Create a horizontal bar chart with the ten largest coefficients from `feat_imp`. Be sure to label your x-axis "Odds Ratio".

- Create a bar chart using pandas.

```
[35]: feat_imp.tail(10).plot(kind="barh")
      plt.xlabel("Odds Ratio");
```




```
[36]: VimeoVideo("665415581", h="d15477e14d", width=600)
```

```
[36]: <IPython.lib.display.VimeoVideo at 0x7f0475a5b1f0>
```

Task 4.4.16: Create a horizontal bar chart with the ten smallest coefficients from `feat_imp`. Be sure to label your x-axis "Odds Ratio".

- Create a bar chart using pandas.

```
[37]: feat_imp.head(10).plot(kind="barh")
      plt.xlabel("Odds Ratio");
```



3.1 Explore Some More

```
[38]: VimeoVideo("665415631", h="90ba264392", width=600)
```

```
[38]: <IPython.lib.display.VimeoVideo at 0x7f0476c51b20>
```

Task 4.4.17: Which municipalities saw the highest proportion of severely damaged buildings? Create a DataFrame `damage_by_vdcmun` by grouping `df` by `"vdcmun_id"` and then calculating the mean of the `"severe_damage"` column. Be sure to sort `damage_by_vdcmun` from highest to lowest proportion.

- Aggregate data using the `groupby` method in `pandas`.

```
[41]: damage_by_vdcmun = (df.groupby("vdcmun_id")["severe_damage"].mean()
    ↪sort_values(ascending=False)).to_frame()
damage_by_vdcmun
```

```
[41]:      severe_damage
vdcmun_id
31      0.930199
32      0.851117
35      0.827145
30      0.824201
33      0.782464
34      0.666979
39      0.572344
40      0.512444
38      0.506425
36      0.503972
```

```
[42]: # Check your work
assert isinstance(
    damage_by_vdcmun, pd.DataFrame
), f"`damage_by_vdcmun` should be a Series, not type {type(damage_by_vdcmun)}."
assert damage_by_vdcmun.shape == (
    11,
    1,
), f"`damage_by_vdcmun` should be shape (11,1), not {damage_by_vdcmun.shape}."
```

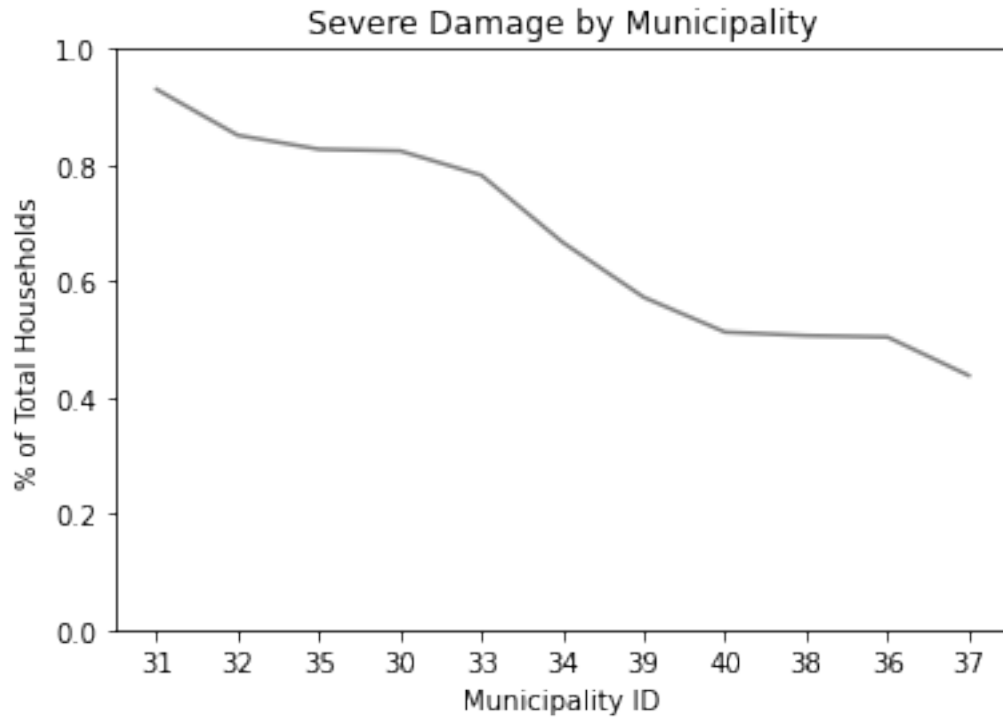
```
[43]: VimeoVideo("665415651", h="9b5244dec1", width=600)
```

```
[43]: <IPython.lib.display.VimeoVideo at 0x7f0471ff7ee0>
```

Task 4.4.18: Create a line plot of `damage_by_vdcmun`. Label your x-axis "Municipality ID", your y-axis "% of Total Households", and give your plot the title "Household Damage by Municipality".

- Create a line plot in Matplotlib.

```
[50]: # Plot line
#damage_by_vdcmun.plot(kind="bar")
plt.plot(damage_by_vdcmun.values, color="grey")
plt.xticks(range(len(damage_by_vdcmun)), labels=damage_by_vdcmun.index)
plt.yticks(np.arange(0.0, 1.1, 0.2))
plt.xlabel("Municipality ID")
plt.ylabel("% of Total Households")
plt.title("Severe Damage by Municipality");
```



Given the plot above, our next question is: How are the Gurung and Kumal populations distributed across these municipalities?

```
[51]: VimeoVideo("665415693", h="fb2e54aa04", width=600)
```

```
[51]: <IPython.lib.display.VimeoVideo at 0x7f047197a070>
```

Task 4.4.19: Create a new column in `damage_by_vdcmun` that contains the the proportion of Gurung households in each municipality.

- [Aggregate data using the groupby method in pandas.](#)
- [Create a Series in pandas.](#)

```
[58]: damage_by_vdcmun["Gurung"] = (
        df[df["caste_household"] == "Gurung"].groupby("vdcmun_id")["severe_damage"].
        ↪count()
        / df.groupby("vdcmun_id")["severe_damage"].count()
    )
    damage_by_vdcmun
```

```
[58]:
```

	severe_damage	Gurung
vdcmun_id		
31	0.930199	0.326937
32	0.851117	0.387849
35	0.827145	0.826889

30	0.824201	0.338152
33	0.782464	0.011943
34	0.666979	0.385084
39	0.572344	0.097971
40	0.512444	0.246727
38	0.506425	0.049023
36	0.503972	0.143178
37	0.437789	0.050485

```
[59]: VimeoVideo("665415707", h="9b29c23434", width=600)
```

```
[59]: <IPython.lib.display.VimeoVideo at 0x7f047197a8e0>
```

Task 4.4.20: Create a new column in `damage_by_vdcmun` that contains the the proportion of Kumal households in each municipality. Replace any NaN values in the column with 0.

- Aggregate data using the `groupby` method in pandas.
- Create a Series in pandas.

```
[61]: damage_by_vdcmun["Kumal"] = (
        df[df["caste_household"] == "Kumal"].groupby("vdcmun_id")["severe_damage"] .
        ↪count()
        / df.groupby("vdcmun_id")["severe_damage"].count()
    ).fillna(0)
    damage_by_vdcmun
```

```
[61]:
```

	severe_damage	Gurung	Kumal
vdcmun_id			
31	0.930199	0.326937	0.000000
32	0.851117	0.387849	0.000000
35	0.827145	0.826889	0.000000
30	0.824201	0.338152	0.000000
33	0.782464	0.011943	0.029478
34	0.666979	0.385084	0.000000
39	0.572344	0.097971	0.000267
40	0.512444	0.246727	0.036973
38	0.506425	0.049023	0.100686
36	0.503972	0.143178	0.003282
37	0.437789	0.050485	0.048842

```
[62]: VimeoVideo("665415729", h="8d0712c306", width=600)
```

```
[62]: <IPython.lib.display.VimeoVideo at 0x7f047197a0a0>
```

Task 4.4.21: Create a visualization that combines the line plot of severely damaged households you made above with a stacked bar chart showing the proportion of Gurung and Kumal households in each district. Label your x-axis "Municipality ID", your y-axis "% of Total Households".

- Create a bar chart using pandas.

- Drop a column from a DataFrame using pandas.

```
[67]: damage_by_vdcmun.drop(columns="severe_damage").plot(
      kind = "bar", stacked=True
    )

plt.plot(damage_by_vdcmun["severe_damage"].values, color="grey")
plt.xticks(range(len(damage_by_vdcmun)), labels=damage_by_vdcmun.index)
plt.yticks(np.arange(0.0,1.1,0.2))
plt.xlabel("Municipality ID")
plt.ylabel("% of Total Households")
plt.title("Severe Damage by Municipality");
```

